



What's new in Spark 2.0?

Rerngvit Yanggratoke @ Combient AB
Örjan Lundberg @ Combient AB

Machine Learning Stockholm Meetup
27 October, 2016
Schibsted Media Group

Combient - Who We Are



- A joint venture, founded in January 2015, owned by Swedish global enterprises
- One of our key missions is the Analytics Centre of Excellence (ACE)
 - Providing *data science and data engineer* resources to our owners and members
 - Sharing best practises, advanced methods, and state-of-the-art technology to generate *business value* from data



SAAB



ERICSSON



Electrolux



LKAB

SKF



SIGMA
Connectivity



stordenso

S|E|B

i investor

FAM

Höganäs

permobil



Aleris



MÖLNLYCKE
HEALTH CARE



Husqvarna®



SÖDRA



WÄRTSILÄ



Spark 2.0 “Easier, Faster, and Smarter”

- Spark Session - a new entry point
- Unification of APIs => Dataset & Dataframe
- Spark SQL enhancements: subqueries, native Window functions, better error messages
- Built-in CSV file type support
- Machine learning pipeline persistence across languages
- Approximate DataFrame Statistics Functions
- Performance enhancement (Whole-stage-code generation)
- Structured streaming (Alpha)



Spark Session

Spark 1.6 => SparkContext, SQLContext, HiveContext

```
val conf = new SparkConf().setMaster(master).setAppName(appName)
val sc    = new SparkContext(conf)
val sqlContext = new org.apache.spark.sql.SQLContext(sc)
val hiveContext = new org.apache.spark.sql.hive.HiveContext(sc)
```

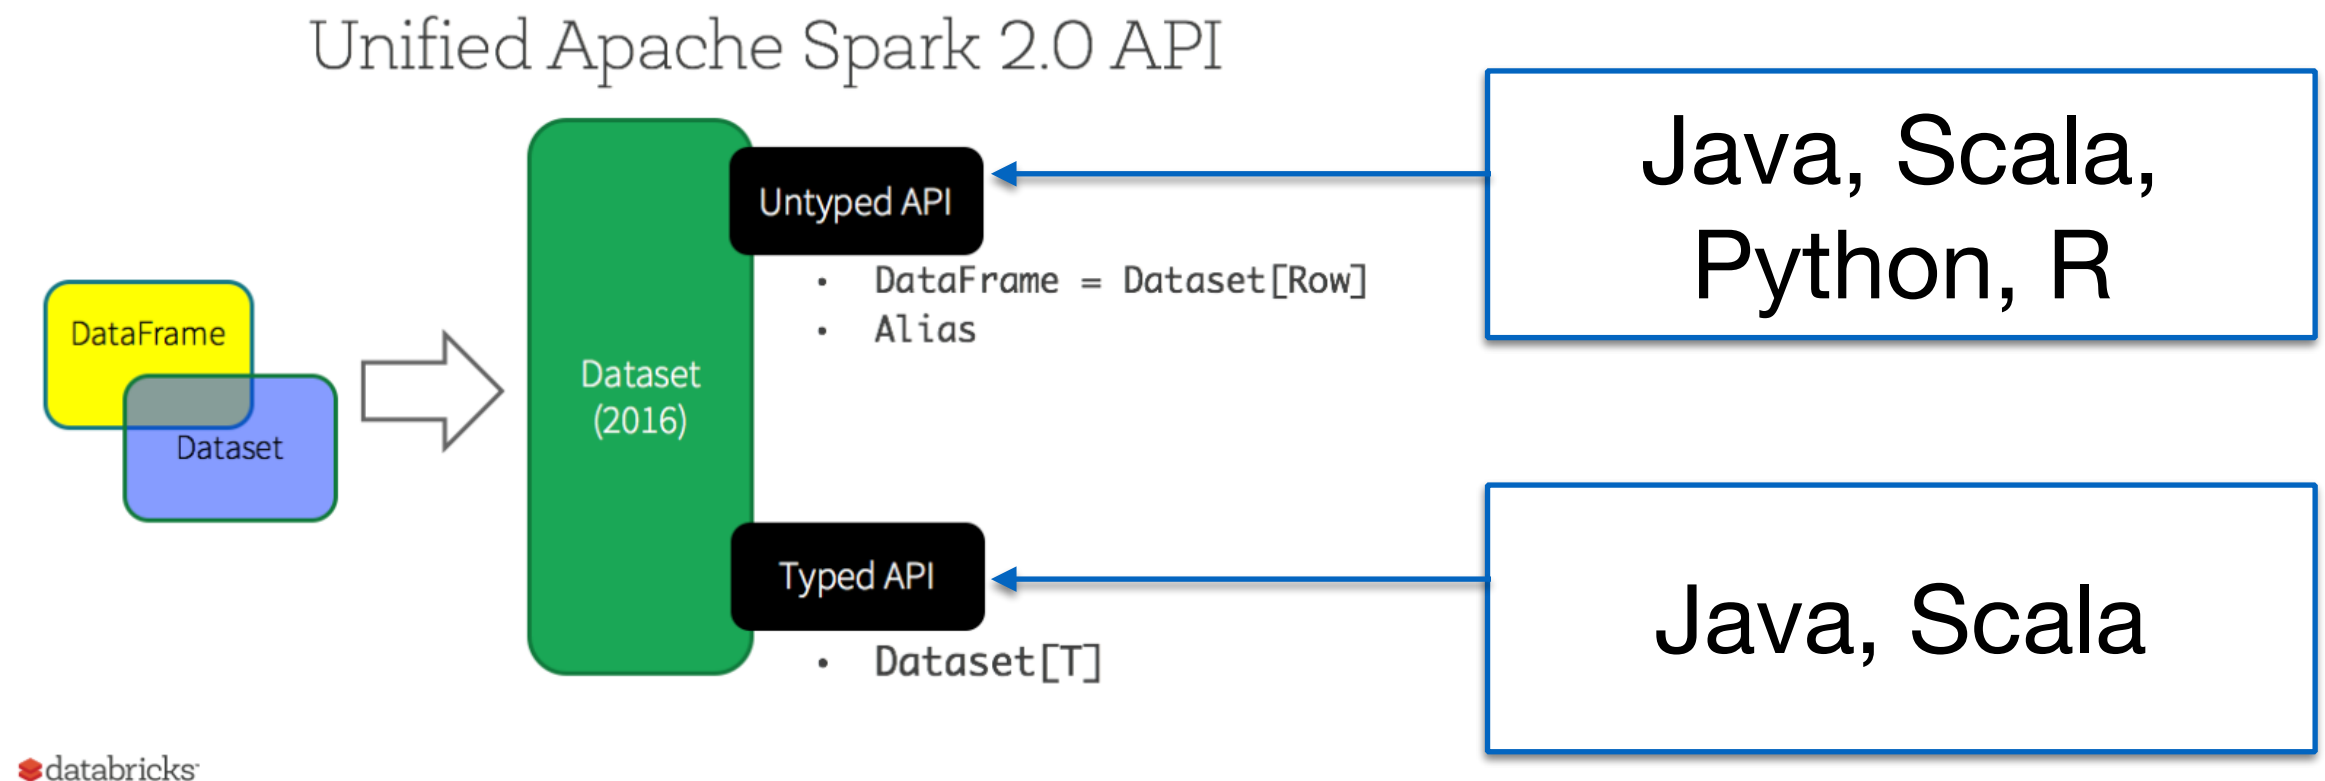
Spark 2.0 => SparkSession

```
val ss = SparkSession.builder().master(master)
    .appName(appName).getOrCreate()
```



Unified API: DataFrame & Dataset

- DataFrame - introduced in Spark 1.3 (March 2015)
- Dataset - experimental in Spark 1.6 (March 2016)



- Observation: very rapid changes of APIs




Dataset API example



<i>name</i>	<i>population_2014</i>	<i>population_2015</i>
<i>Afghanistan</i>	<i>XXXX</i>	<i>XXXX</i>
<i>Albania</i>	<i>XXXX</i>	<i>XXXX</i>
<i>....</i>	<i>....</i>	<i>....</i>

Table: country_pop



What countries has a shrinking population from 2014 to 2015?

```
> case class CountryData(name: String, population_2014: Long, population_2015: Long)
val countryDS = ss.sql(" select * from country_pop ").as[CountryData]
countryDS.filter(c => {c.population_2015 < c.population_2014}).show(10, false)
```

```
+-----+-----+-----+
|name    |population_2014|population_2015|
+-----+-----+-----+
|Japan   |127131800      |126958472      |
|Spain   |46480882       |46418269       |
|Ukraine |45362900       |45198200       |
|Poland  |38011735       |37999494       |
|Romania |19908979       |19832389       |
+-----+-----+-----+
```

Compile time
error detection

- Key benefit over DataFrame API: static typing!



Subqueries

Ex. Which countries have more than 100M people in 2016?

Spark 1.6 => Subqueries inside FROM clause

```
> ss.sql(""" SELECT *  
              FROM (SELECT * FROM country_pop WHERE population_2014 > 100e6) t1  
              ORDER BY population_2014 DESC LIMIT 5""").show(10, false)
```

Spark 2.0 => Subqueries inside FROM and WHERE clause

```
> ss.sql(""" SELECT *  
              FROM country_pop  
              WHERE population_2014 IN (select population_2014 from country_pop WHERE population_2014 > 100e6)  
              ORDER BY NAME ASC LIMIT 5 """).show(10, false)
```

name	population_2014	population_2015
Bangladesh	159077513	160995642
Brazil	206077898	207847528
China	1364270000	1371220000
India	1295291543	1311050527
Indonesia	254454778	257563815



Window Functions

Spark 1.X => required HiveContext

Spark 2.0 => native window function

```
> // PARTITION BY id ORDER BY cykle ROWS BETWEEN 2 PRECEDING AND 2 FOLLOWING (5)
import org.apache.spark.sql.expressions.Window
import org.apache.spark.sql.functions._
val w = Window.partitionBy("id").orderBy("cykle").rowsBetween(-2, 2)

val x = dft.select($"id", $"cykle", avg($"value").over(w))
x.show
```

```
dft.createOrReplaceTempView("dft")
val sqlq = """
    SELECT id,cykle,avg(value) OVER (PARTITION BY id ORDER BY cykle ASC ROWS BETWEEN 2 PRECEDING AND 2 FOLLOWING)
    FROM dft
"""

val dfsQL = sqlContext.sql(sqlq)
```




SQL Error messages

Spark 1.6 => an exception happen at this location X

```
> val sqlq = """
    SELECT id,cykle,avg(value) OVERT (PARTITION BY id ORDER BY cykle ASC ROWS BETWEEN 2 PRECEDING AND 2 FOLLOWING)
    FROM dft
    """

val dfsQL = sqlContext.sql(sqlq)
```

org.apache.spark.sql.AnalysisException: missing EOF at '(' near 'OVERT'; line 1 pos 33

```
at org.apache.spark.sql.hive.HiveQL$.createPlan(HiveQL.scala:318)
at org.apache.spark.sql.hive.ExtendedHiveQLParser$$anonfun$hiveQL$1.apply(ExtendedHiveQLParser.scala:41)
at org.apache.spark.sql.hive.ExtendedHiveQLParser$$anonfun$hiveQL$1.apply(ExtendedHiveQLParser.scala:40)
at scala.util.parsing.combinator.Parsers$Success.map(Parsers.scala:136)
at scala.util.parsing.combinator.Parsers$Success.map(Parsers.scala:135)
```

Spark 2.0 => illustrate possible options

```
> val sqlq = """
    SELECT id,cykle,avg(value) OVERT (PARTITION BY id ORDER BY cykle ASC ROWS BETWEEN 2 PRECEDING AND 2 FOLLOWING)
    FROM dft
    """

val dfsQL = sqlContext.sql(sqlq)
```

org.apache.spark.sql.catalyst.parser.ParseException:

extraneous input '(' expecting {<EOF>, ',', 'FROM', 'WHERE', 'GROUP', 'ORDER', 'HAVING', 'LIMIT', 'LATERAL', 'WINDOW', 'UNION', 'EXCEPT', 'INTERSECT', 'SORT', 'CLUSTER', 'DISTRIBUTE'}(line 2, pos 45)

== SQL ==

```
SELECT id,cykle,avg(value) OVERT (PARTITION BY id ORDER BY cykle ASC ROWS BETWEEN 2 PRECEDING AND 2 FOLLOWING)
```

```
-----^^^
FROM dft
```



Built-in CSV file support

Spark 1.6 => require package databricks/spark-csv


```
> val df = sqlContext.read.format("com.databricks.spark.csv").option("header", "true")  
  .load("s3n://com-combient-test/ml-meetup/The_world_population_data.csv")
```

Spark 2.0 => built-in

```
> val df = spark.read.option("header", "true")  
  .csv("s3n://com-combient-test/ml-meetup/The_world_population_data.csv")
```

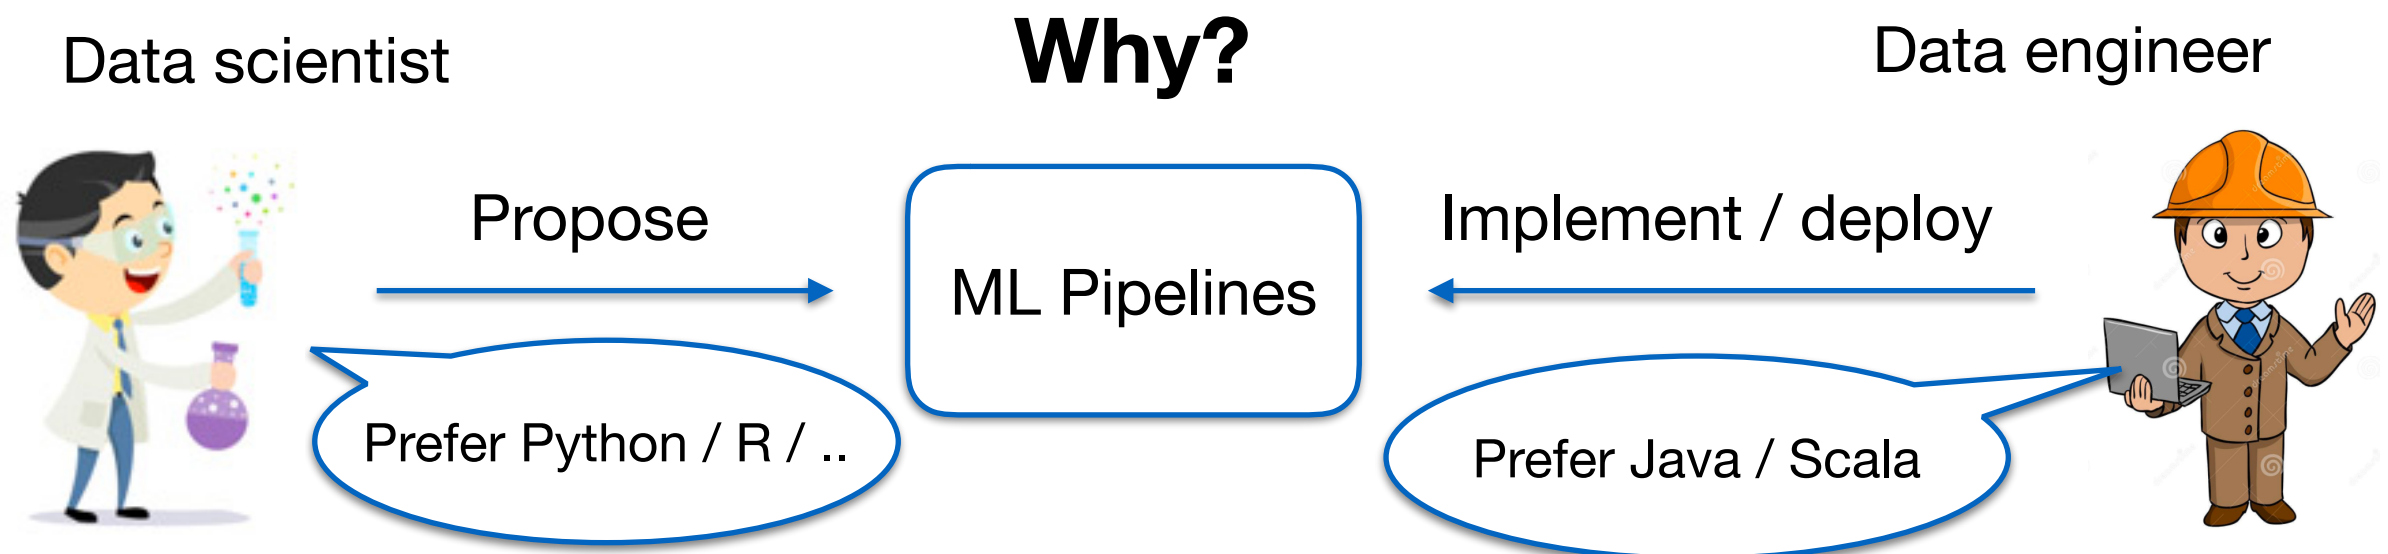
Spark 2.0 Machine learning (ML) APIs Status

- spark.mllib (RDD-based APIs)
 - Entering bug-fixing modes (no new features)
- spark.ml (Dataframe-based): models + transformations
 - Introduced in January 2015
 - Become a primary API



An example Pipeline model

ML Pipeline Persistence Across languages



Source: classroomclipart.com

Source: worldartsme.com/

- Prototype and experiments with candidate ML pipelines
- Focus on rapid prototype and fast exploration of data

- Implement and deploy pipelines in production environment
- Focus on maintainable and reliable systems

Pipeline persistence allow for **import/export** pipeline across languages

- Introduced only for Scala in Spark 1.6
- Status of Spark 2.0
 - Scala/Java (Full), Python (Near-full), R (with hacks)



Create ML pipeline and persist it in Python

```
> %py
# Building a pipeline
from pyspark.ml import Pipeline
from pyspark.ml.feature import StandardScaler
from pyspark.ml.classification import RandomForestClassifier
scaler = StandardScaler(inputCol="features", outputCol="scaledFeatures",
                        withStd=True, withMean=False)
scaler_model = scaler.fit(training_df)
dataset_scaled = scaler_model.transform(training_df)
rf_classifier = RandomForestClassifier(featuresCol="scaledFeatures", labelCol="label",
                                     predictionCol="prediction", probabilityCol="probability",
                                     maxDepth=5, numTrees=100)
ml_pipeline = Pipeline(stages=[scaler_model, rf_classifier])

# Fitting the pipeline model to the training set
pipeline_model = ml_pipeline.fit(training_df)

# Writing the pipeline to a file
base_path = "/tmp/mlpipe-persistence-example"
pipeline_model.save(base_path + "/fitted_pipeline")
```




Loading the pipeline and test it in Scala

```
> import org.apache.spark.ml._
val base_path = "/tmp/mlpipe-persistence-example"
// Loading the pipeline
val pipeline = PipelineModel.read.load(base_path + "/fitted_pipeline")

// Testing the pipeline
val predictions = pipeline.transform(test_df)
predictions.show
```

label	features	scaledFeatures	rawPrediction	probability prediction
7.0	(784,[202,203,204...	(784,[202,203,204...	[2.42187287483487...	[0.02421872874834... 7.0
2.0	(784,[94,95,96,97...	(784,[94,95,96,97...	[9.73366907185490...	[0.09733669071854... 2.0
1.0	(784,[128,129,130...	(784,[128,129,130...	[0.28122374573791...	[0.00281223745737... 1.0
0.0	(784,[124,125,126...	(784,[124,125,126...	[77.5100579522886...	[0.77510057952288... 0.0
4.0	(784,[150,151,159...	(784,[150,151,159...	[3.55187312690512...	[0.03551873126905... 4.0
1.0	(784,[156,157,158...	(784,[156,157,158...	[0.18799849144322...	[0.00187998491443... 1.0
4.0	(784,[149,150,151...	(784,[149,150,151...	[1.78116677229671...	[0.01781166772296... 4.0
9.0	(784,[179,180,181...	(784,[179,180,181...	[0.95908360963345...	[0.00959083609633... 9.0
5.0	(784,[129,130,131...	(784,[129,130,131...	[6.89697832373067...	[0.06896978323730... 6.0
9.0	(784,[209,210,211...	(784,[209,210,211...	[1.72278921501600...	[0.01722789215016... 9.0
0.0	(784,[123,124,125...	(784,[123,124,125...	[88.0199490767903...	[0.88019949076790... 0.0
6.0	(784,[94,95,96,97...	(784,[94,95,96,97...	[15.5591258414737...	[0.15559125841473... 6.0
9.0	(784,[208,209,210...	(784,[208,209,210...	[1.58963036387379...	[0.01589630363873... 9.0
0.0	(784,[152,153,154...	(784,[152,153,154...	[76.4044263510522...	[0.76404426351052... 0.0
1.0	(784,[125,126,127...	(784,[125,126,127...	[0.21724594335647...	[0.00217245943356... 1.0
5.0	(784,[124,125,126...	(784,[124,125,126...	[10.5215301036557...	[0.10521530103655... 3.0
9.0	(784,[179,180,181...	(784,[179,180,181...	[2.80678452853574...	[0.02806784528535... 9.0
7.0	(784,[200,201,202...	(784,[200,201,202...	[3.93447536107229...	[0.03934475361072... 7.0
3.0	(784,[118,119,120...	(784,[118,119,120...	[3.49591652293212...	[0.03495916522932... 6.0
4.0	(784,[158,159,185...	(784,[158,159,185...	[1.98901146497452...	[0.01989011464974... 4.0

Approximate DataFrame Statistics Functions



Motivation: rapid explorative data analysis

Bloom filter - Approximate set membership data structure

Burton H. Bloom. 1970. Space/time trade-offs in hash coding with allowable errors. Commun. ACM 13, 7 (July 1970), 422-426. DOI=<http://dx.doi.org/10.1145/362686.362692>

CountMinSketch - Approximate frequency estimation data structure

Cormode, Graham (2009). "Count-min sketch" (PDF). Encyclopedia of Database Systems. Springer. pp. 511–516.


Approximate Quantile

Michael Greenwald and Sanjeev Khanna. 2001. Space-efficient online computation of quantile summaries. In Proceedings of the 2001 ACM SIGMOD international conference on Management of data (SIGMOD '01), Timos Sellis and Sharad Mehrotra (Eds.). ACM, New York, NY, USA, 58-66.

Bloom filter- Approximate set membership data structure



Inserting items $O(k)$



k hash functions

Source : wikipedia

w

Membership query $\Rightarrow O(k)$

Trade space efficiency with false positives

Spark 2: example code

```
> val bfCountryName = df.stat.bloomFilter(colName="name", expectedNumItems=300, fpp =0.01)
println(bfCountryName.mightContain("Sweden"))
println(bfCountryName.mightContain("Unknow Country"))
```

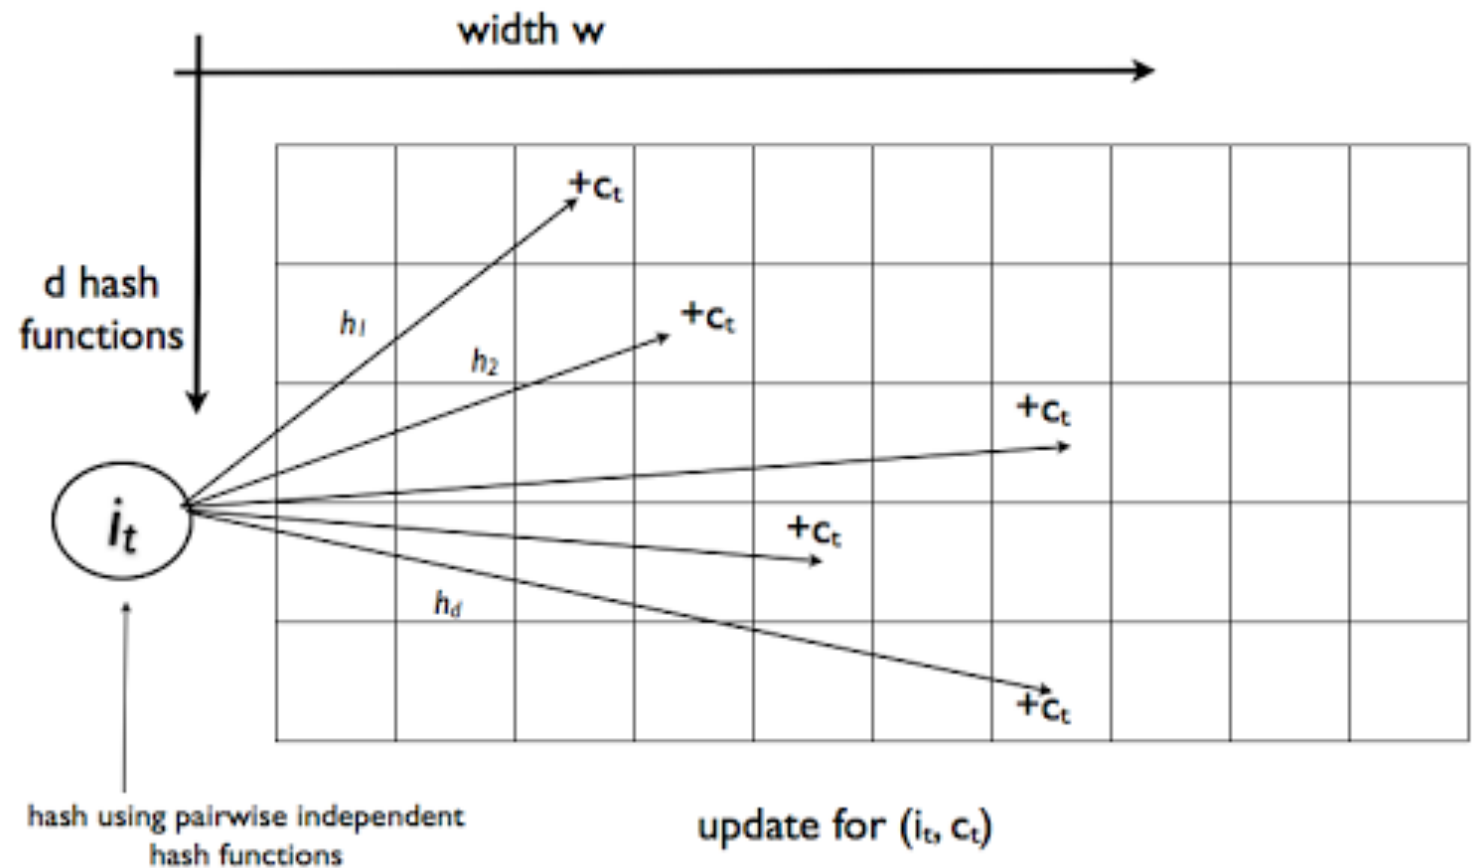
true
false



Count-Min Sketch

Operations

- Add item - $O(wd)$
- Query item frequency - $O(wd)$



Source : <http://debasishg.blogspot.se/2014/01/count-min-sketch-data-structure-for.html>

Spark 2: example code

```
> val nameSketch = df.stat.countMinSketch(colName="name", eps = 0.001,
      confidence = 0.99, seed = 42)
println(nameSketch.estimateCount("Sweden"))
println(nameSketch.estimateCount("Unknown Country"))
```

1

0



Approximate Quantile

Details are much more complicated than the above two. :)

Spark 2: example code

```
> df.stat.approxQuantile(  
    col="population_2014",  
    probabilities=Array(0.25, 0.5, 0.75, 0.9),  
    relativeError=0.05)
```



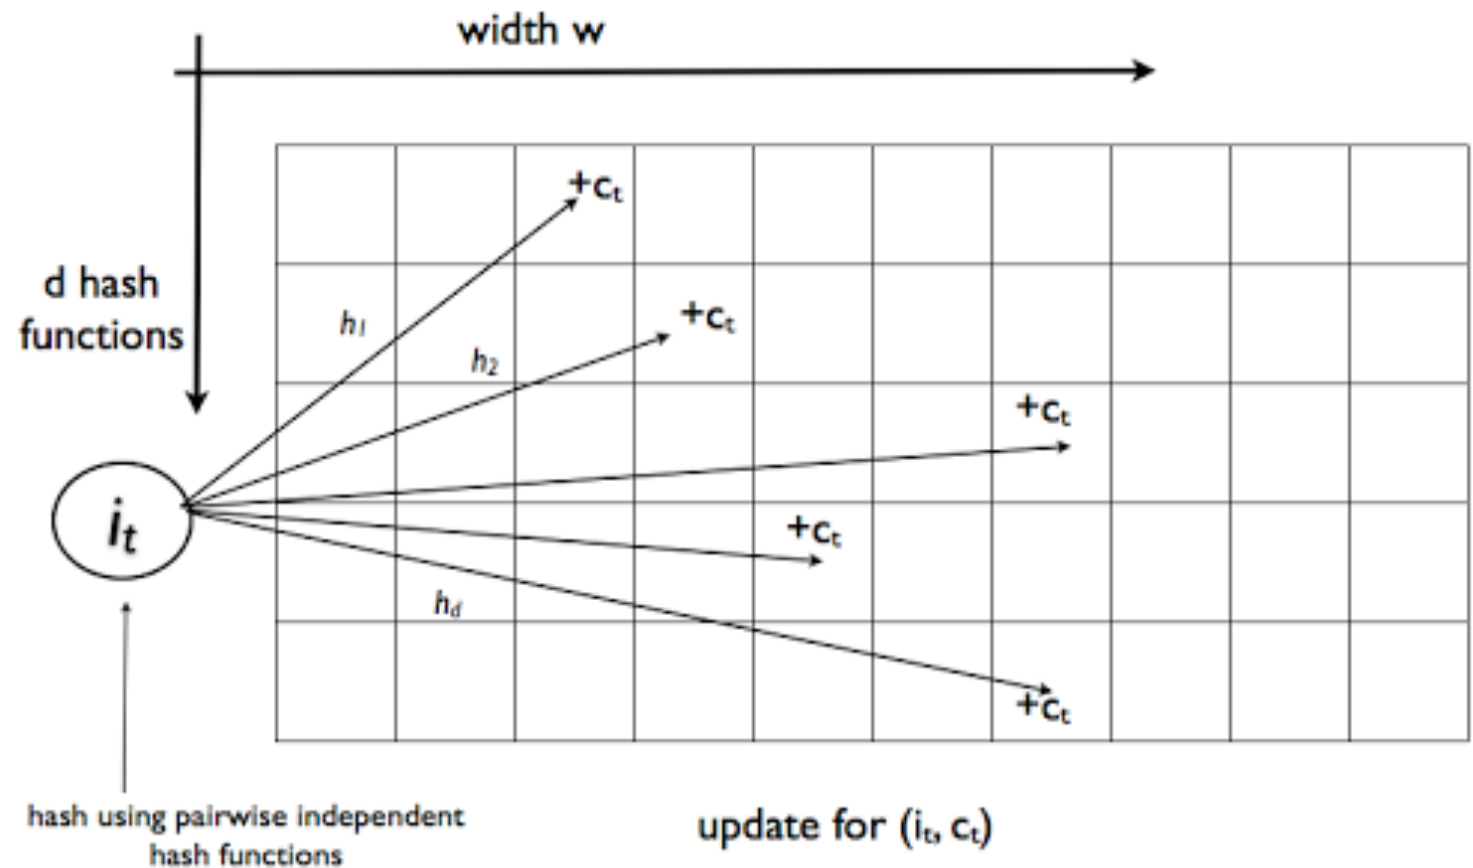
```
res189: Array[Double] = Array(2219937.0, 1.0598482E7, 1.9113728E7, 2.6786598E7)
```



Count-Min Sketch

Operations

- Add item - $O(wd)$
- Query item frequency - $O(wd)$



Source : <http://debasishg.blogspot.se/2014/01/count-min-sketch-data-structure-for.html>

Spark 2: example code

```
> val nameSketch = df.stat.countMinSketch(colName="name", eps = 0.001,
                                           confidence = 0.99, seed = 42)
println(nameSketch.estimateCount("Sweden"))
println(nameSketch.estimateCount("Unknown Country"))
```

1

0



Approximate Quantile

Details are much more complicated than the above two. :)

Spark 2: example code

```
> df.stat.approxQuantile(  
  col="population_2014",  
  probabilities=Array(0.25, 0.5, 0.75, 0.9),  
  relativeError=0.05)
```



res8: Array[Double] = Array(167543.0, 1260934.0, 3753121.0, 5643475.0)

Databrick example notebook:

https://docs.cloud.databricks.com/docs/latest/sample_applications/04%20Apache%20Spark%202.0%20Examples/05%20Approximate%20Quantile.html



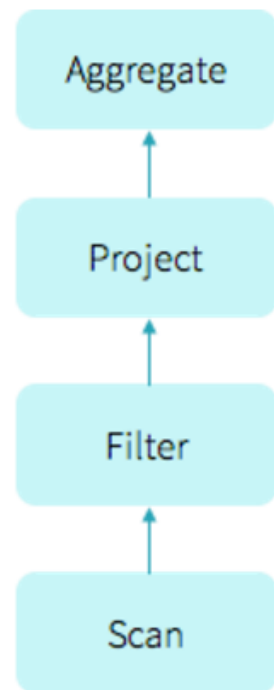
Whole-stage code generation

- Collapsing multiple function calls into one
 - Minimising overheads (e.g., virtual function calls) and using CPU registers



Example

```
select count(*) from store_sales
where ss_item_sk = 1000
```

Source:  databricks



Volcano (Spark 1.X)



```
var count = 0
for (ss_item_sk in store_sales) {
  if (ss_item_sk == 1000) {
    count += 1
  }
}
```

Whole-stage code generation
(Spark 2.0)

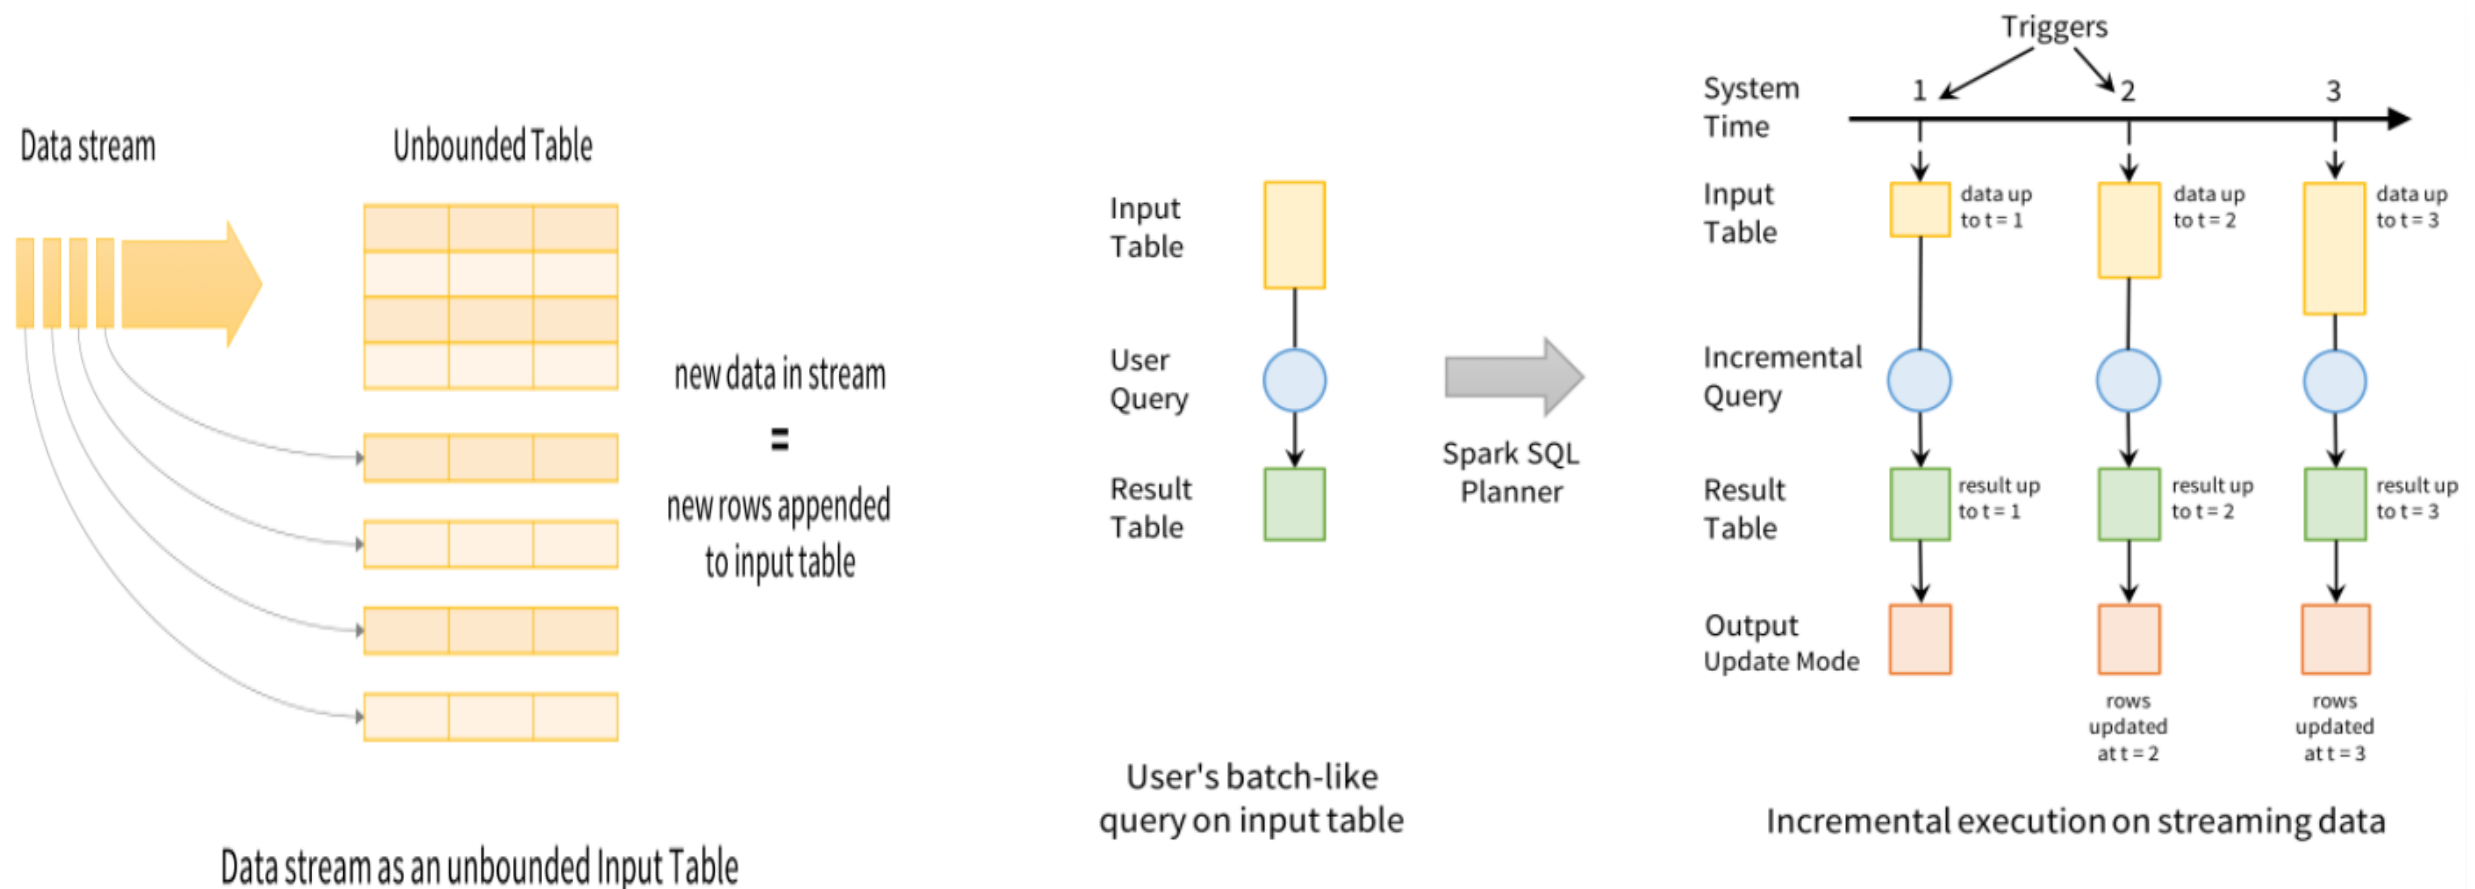
Databrick example notebook:

<https://databricks-prod-cloudfront.cloud.databricks.com/public/4027ec902e239c93eaaa8714f173bcfc/6122906529858466/293651311471490/5382278320999420/latest.html>



Structure Streaming [Alpha]

- SQL on streaming data



Source : <https://databricks.com/blog/2016/07/28/structured-streaming-in-apache-spark.html>

Databrick example notebook:

<https://databricks-prod-cloudfront.cloud.databricks.com/public/4027ec902e239c93eaaa8714f173bcfc/4012078893478893/295447656425301/5985939988045659/latest.html>



Structure Streaming [Alpha]

Batch Version

```
// Read JSON once from S3
logsDF = spark.read.json("s3://logs")

// Transform with DataFrame API and save
logsDF.select("user", "url", "date")
        .write.parquet("s3://out")
```

Streaming Version

Not ready for production

```
// Read JSON continuously from S3
logsDF = spark.readStream.json("s3://logs")

// Transform with DataFrame API and save
logsDF.select("user", "url", "date")
        .writeStream.parquet("s3://out")
        .start()
```

Source: <https://databricks.com/blog/2016/07/28/continuous-applications-evolving-streaming-in-apache-spark-2-0.html>



We are currently recruiting 😊
combient.com/#careers